

User Manual of MLP simulator (+NeuroSim)

Developers: Pai-Yu Chen and Xiaochen Peng

PI: Prof. Shimeng Yu, Arizona State University

Index

1. Introduction.....	1
2. System Requirements (Linux)	2
3. Installation and Usage (Linux).....	2
4. Device Level: Extraction of Synaptic Device Characteristics	3
4.1 Non-ideal Analog eNVM Device Properties.....	3
4.2 Fitting by MATLAB script (nonlinear_fit.m)	4
5. Circuit Level: Synaptic Cores and Array Architectures	7
5.1 Crossbar Array Architecture.....	7
5.2 Pseudo-crossbar Array Architecture.....	8
5.3 Array Peripheral Circuits.....	9
6. Algorithm Level: Multilayer Perceptron (MLP) Neural Network Architecture	12
7. How to run MLP simulator (+NeuroSim).....	14

1. Introduction

MLP simulator (+NeuroSim) is developed in C++ to emulate the online learning/offline classification scenario with MNIST handwritten dataset in a 2-layer multilayer perceptron (MLP) neural network based on emerging non-volatile memory (eNVM) array architectures. The eNVM in this simulator refers to a special subset of resistive memory devices that can tune the conductance into multilevel states with voltage stimulus. *NeuroSim* is a circuit-level macro model for benchmarking neuro-inspired architectures in terms of circuit-level performance metrics, such as chip area, latency, dynamic energy and leakage power. Without *NeuroSim*, *MLP simulator* can be regarded as a standalone functional simulator that is able to evaluate the learning accuracy and the circuit-level performance (but only for the synapse array) during learning. With *NeuroSim*, *MLP simulator (+NeuroSim)* becomes an integrated framework with hierarchical organization from the device level (transistor and analog eNVM device properties) to the circuit level (analog eNVM array architectures with periphery circuit modules) and then to the algorithm level (neural network topology), enabling trace-based and cycle-accurate evaluation on the learning accuracy as well as the circuit-level performance metrics at the run-time of learning.

In this released version 1.0, the target users are device engineers who wish to quickly estimate the system-level performance with his/her own analog synaptic device data. The users are expected to have the weight

update characteristics (conductance vs. # pulse) ready in hand. Device-level parameters such as number of levels, weight update nonlinearity, device-to-device variation, and cycle-to-cycle variations, could be extracted using the MATLAB script that is provided. At the circuit level, a few design options are available, such as the array architecture being true crossbar or pseudo crossbar with 1T1R. At the algorithm level, a simple 2-layer MLP neural network is provided for evaluation, thus only limited options are available to the users to modify, such as the size of each layer and the size of weight matrices.

2. System Requirements (Linux)

The tool is expected to run in Linux with required system dependencies installed. These include GCC, GNU make, GNU C libraries (glibc). We have tested the compatibility of the tool with a few different Linux environments, such as (1) Red Hat 5.11 (Tikanga), gcc v4.7.2, glibc 2.5, (2) Red Hat 7.3 (Maipo), gcc v4.8.5, glibc v2.1.7, (3) Ubuntu 16.04, gcc v5.4.0, glibc v2.23, and they are all workable.

※ The tool may not run correctly (stuck forever) if compiled with gcc 4.5 or below, because some C++11 features are not well supported.

3. Installation and Usage (Linux)

Step 1: Get the tool from GitHub

```
git clone https://github.com/neurosim/MLP_NeuroSim.git
```

Step 2: Extract **MNIST_data.zip** to it's current directory

```
unzip MNIST_data.zip
```

Step 3: Compile the codes

```
make
```

Summary of the useful commands is provided below. It is recommended to execute these commands under the tool's directory.

Command	Description
make	Compile the codes and build the “main” program
make clean	Clean up the directory by removing the object files and the “main” executable
./main	Run simulation (after make)
make run	Run simulation (after make), and the results will be saved to a log file (filename appended with the current time info). This command does not work if “stdbuf” is not found.

※ The simulation uses OpenMP for multithreading, and it will use up all the CPU cores by default.

4. Device Level: Extraction of Synaptic Device Characteristics

4.1 Non-ideal Analog eNVM Device Properties

As shown in Fig. 1, the framework considers the following non-ideal analog eNVM device properties:

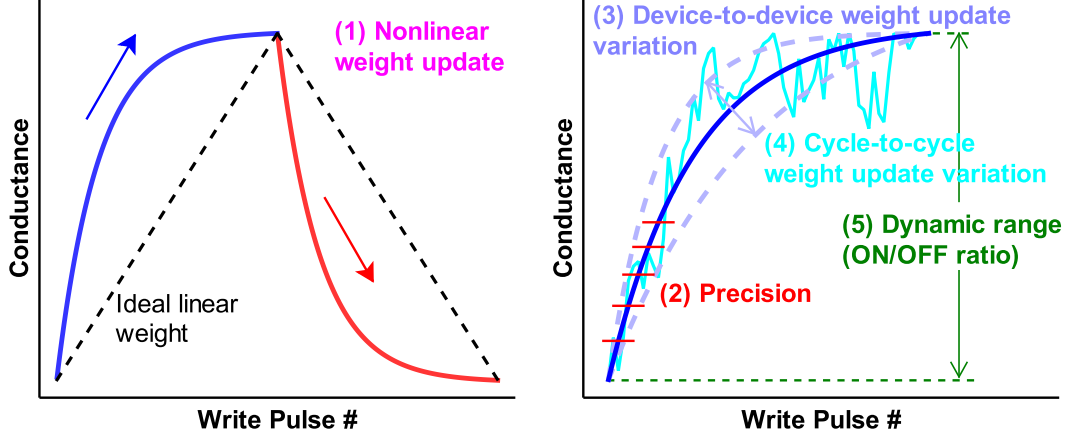


Fig. 1 Summary of non-ideal analog eNVM device properties.

1) Nonlinear weight update

Ideally, the amount of weight increase (or long-term potentiation, LTP) and weight decrease (or long-term depression, LTD) should be linearly proportional to the number of write pulses. However, the realistic devices reported in literature do not follow such ideal trajectory, where the conductance typically changes rapidly at the beginning stages of LTP and LTD and then gradually saturates. We have built a device behavioral model to capture nonlinear weight update behavior, where the conductance change with number of pulses (P) is described with the following equations:

$$G_{LTP} = B \left(1 - e^{-\left(\frac{P}{A}\right)} \right) + G_{min} \quad (1)$$

$$G_{LTD} = -B \left(1 - e^{-\left(\frac{P - P_{max}}{A}\right)} \right) + G_{max} \quad (2)$$

$$B = (G_{max} - G_{min}) / (1 - e^{-\frac{P_{max}}{A}}) \quad (3)$$

G_{LTP} and G_{LTD} are the conductance for LTP and LTD, respectively. G_{max} , G_{min} , and P_{max} are directly extracted from the experimental data, which represents the maximum conductance, minimum conductance and the maximum pulse number required to switch the device between the minimum and maximum conductance states. A is the parameter that controls the nonlinear behavior of weight update. A can be positive (blue) or negative (red). In Fig. 1, the A of LTP and LTD has the same magnitude but different signs. B is simply a function of A that fits the functions within the range of G_{max} , G_{min} , and P_{max} . All these parameters can be different in LTP and LTD in the fitting by the MATLAB script. However, for simplicity, the simulator currently uses the smaller value of G_{max} and the larger value of G_{min} in LTP and LTD.

Using Eq. (1)-(3), a set of nonlinear weight increase (blue) and weight decrease (red) behavior can be obtained by adjusting A as shown in Fig. 2, where each nonlinear curve is labeled with a nonlinearity value from +6 to -6. It can be proved that Eq. (1) and (2) are equivalent with a different sign of A , thus we will just use Eq. (1) to calculate both nonlinear LTP and LTD weight update. Different than Fig. 1, all LTD

curves are mirrored and shifted horizontally to make sure the curve starting from the pulse number 0 for simpler formulization.




Fig. 2 Analog eNVM device behavioral model of the nonlinear weight update with the nonlinearity labeled from -6 to 6.

2) Limited precision

The precision of an analog eNVM device is determined by the number of conductance states it has, which is P_{\max} in Eq. (1)-(3).

3) Device-to-device weight update variation

The effect of device-to-device weight update variation can be analyzed by introducing the variation into the nonlinearity baseline. This variation is defined as the nonlinearity baseline's standard deviation (σ) respect to 1 step of the 6 steps in Fig. 2.

4) Cycle-to-cycle weight update variation

The Cycle-to-cycle weight update variation is referred to as the variation in conductance change at every programming pulse. This variation (σ) is expressed in terms of the percentage of entire conductance range.

5) Dynamic range (ON/OFF ratio)

Ideally, the weight values are represented by a normalized conductance of analog eNVM devices with the range from 0 to 1. However, the minimum conductance can be regarded as 0 only when the ratio between the maximum and minimum conductance (ON/OFF ratio) approaches infinity. With limited ON/OFF ratio, the cells with weight=0 still have leakage.

4.2 Fitting by MATLAB script (nonlinear_fit.m)

In this section, we will fit the experimental weight update data and extract the device parameters that will be used in the simulator. We have developed a MATLAB script **nonlinear_fit.m** to do such a task, where it has been set up for fitting Ag:a-Si devices in the following reference as an example.

S. H. Jo, T. Chang, I. Ebong, B. B. Bhadviya, P. Mazumder, and W. Lu, "Nanoscale memristor device as synapse in neuromorphic systems," *Nano Lett.*, vol. 10, no. 4, pp. 1297–1301, 2010.

Before the fitting, the user has to make sure the experimental weight update data are pre-processed in a format that is similar to Fig. 2. Namely, the LTD data should be mirrored horizontally and both LTP and LTD data should start from the pulse number 0 so that the data can be fit by Eq. (1)-(3). The user can look at the pre-processed data of Ag:a-Si devices as an example in the MATLAB script.

The shape of nonlinear weight update curves can look very different with the same A but different P_{\max} and G_{\min} , because A has to be scaled according to different P_{\max} . First, we normalize P_{\max} to be 1 by default definition, then we can tune the normalized A and the cycle-to-cycle weight update variation for both LTP and LTD in the generated **Figure 1** (normalized conductance vs. normalized number of pulses) to find the best fit, as shown in Fig. 3. A good procedure is that the user first finds out a reasonable normalized A for LTP and LTD curves without variation (by setting the variation in LTP and LTD to zero), and then try to fit the LTP and LTD data with good variation values and pseudorandom seeds (for example, **rng(103)** and **rng(898)** in script). In the script, the A values are defined as **A_LTP** and **A_LTD**, and P_{\max} is defined as **xf** which is set to be 1. These parameters are shown in Fig. 4. It should be noted that the device-to-device weight update variation is not considered in this script.




Fig. 3 Fitting of Ag:a-Si weight update data with normalized A in the plot of normalized conductance vs. normalized number of pulses.

```

192 - xf = 1;
193 - A_LTP = 0.5;
194 - B_LTP = 1./(1-exp(-1./A_LTP));
195 - A_LTD = -0.2;
196 - B_LTD = 1./(1-exp(-1./A_LTD));

198 % LTP fitting
199 - var_amp = 0.035; % LTP cycle-to-cycle variation
200 - rng(103);

217 % LTD fitting
218 - var_amp = 0.025; % LTD cycle-to-cycle variation
219 - rng(898);

```

Fig. 4 Code snippet of parameters in **nonlinear_fit.m**

After the fitting is done, the user can look up the magnitude of the normalized A value to figure out its corresponding nonlinearity label. We provide the information of one-to-one mapping of the nonlinearity label values to the normalized A values in the file **Nonlinearity-NormA.htm**. The nonlinearity label value is from 0 to 7 with a step of 0.01, which is precise enough. If the normalized A value is negative, the nonlinearity label value will also simply be negative. In the example of Ag:a-Si, the fitted normalized A values are 0.5 and -0.2, and we can figure out their corresponding nonlinearities are 2.40 and -4.88 for LTP and LTD, respectively, as shown in Table 1. In the next section, *MLP simulator (+NeuroSim)* will take these two values as eNVM cell parameters.

Nonlinearity	Norm. A	Nonlinearity	Norm. A	Nonlinearity	Norm. A	Nonlinearity	Norm. A	Nonlinearity	Norm. A
2.31	0.5207	2.38	0.5038	2.45	0.4879	4.84	0.2030	4.91	0.1983
2.32	0.5183	2.39	0.5015	2.46	0.4856	4.85	0.2023	4.92	0.1976
2.33	0.5158	2.40	0.4992	4.86	0.2016	4.93	0.1970
2.34	0.5134	2.41	0.4969	4.87	0.2010	4.94	0.1963
2.35	0.5110	2.42	0.4946	4.88	0.2003	4.95	0.1957
2.36	0.5086	2.43	0.4932	4.82	0.2044	4.89	0.1996	4.96	0.1950
2.37	0.5062	2.44	0.4901	4.83	0.2037	4.90	0.1990	4.97	0.1944

Table 1 Snippet of **Nonlinearity-NormA.htm**

5. Circuit Level: Synaptic Cores and Array Architectures

In this framework, we consider two synaptic cores of 2-layer MLP in the evaluation for circuit-level metrics such as area, latency, energy, etc. A synaptic core is a computation unit that is specifically designed for weighted sum and weight update. It consists of the synaptic array and array periphery. In the simulator, a synaptic core can be instantiated from **SubArray** class in **SubArray.cpp**. In this released version, there are two available design options for the synaptic core. One is based on the crossbar array architecture, and other one is based on the pseudo-crossbar array architecture, as shown in Fig. 5. The details of both architectures and their peripheral circuits are introduced below.

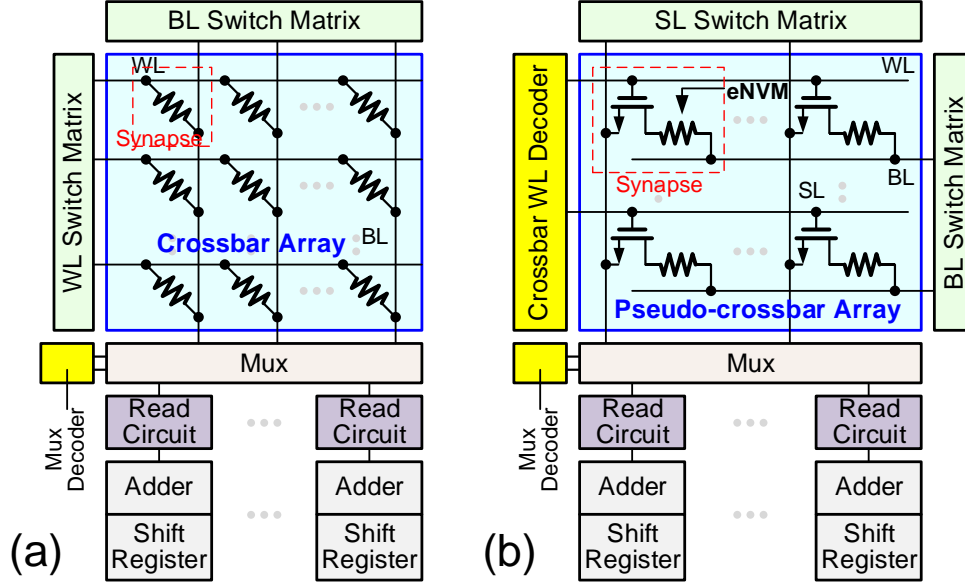


Fig. 5 Synaptic cores based on (a) crossbar and (b) pseudo-crossbar array architectures

5.1 Crossbar Array Architecture

The crossbar array structure has the most compact and simplest array structure for eNVM devices to form a weight matrix, where each eNVM device is located at the cross point of a word line (WL) and a bit line (BL). The crossbar array structure can achieve a high integration density of $4F^2/\text{cell}$ (F is the lithography feature size). If the input vector is encoded by read voltage signals, the weighted sum operation (matrix-vector multiplication) can be performed in a parallel fashion with the crossbar array. However, as there is no isolation between cells, it is necessary to apply some voltage (smaller than the programming voltage, e.g. $V/2$) at all the unselected rows and columns to prevent the write disturbance on unselected cells during weight update. The voltage bias schemes for weight update is shown in Fig. 6. In addition, a two-terminal threshold switching selector device is desired to minimize the write disturbance and sneak path problem. In our model, a simple I-V nonlinearity (**NL**) in **Cell.cpp** is used to define the effect of selector or built-in self-selection. As the weight increase and decrease need different programming voltage polarities, the weight update process requires 2 steps with different voltage bias schemes. In weight update, the selected cells will be on the same row, and programming pulses or biases (if no update) are provided from the BL, allowing the selected cells to be tuned differently in parallel. To perform weight update for the entire array, a row-by-row operation is necessary. Typically, the entire row will be selected at a time to ensure the maximum parallelism. In the simulator, the crossbar array architecture can be designated by setting **cmosAccess=false** in **RealDevice** class of **Cell.cpp**.




Fig. 6 Voltage bias scheme in the write operation of crossbar array. Two separate phases for weight increase and decrease are required. In this example, the left cell of the selected cells will be updated in phase 1, while the right one will be updated in phase 2.

5.2 Pseudo-crossbar Array Architecture

Another common design solution to the write disturbance and sneak path problem is to add a cell selection transistor in series with the eNVM device, forming the one-transistor one-resistor (1T1R) array architecture, as shown in Fig. 7(a). The WL controls the gate of the transistor, which can be viewed as a switch for the cell. The source line (SL) connects to the source of the transistor. The eNVM cell's top electrode connects to the BL, while its bottom electrode connects to the drain of the transistor through a contact via. In such case, the cell area of 1T1R array is then determined by the transistor size, which is typically $>6F^2$ depending on the maximum current required to be delivered into the eNVM cell. Larger current needs larger transistor gate width/length (W/L). However, conventional 1T1R array is not able to perform the parallel weighted sum operation. To solve this problem, we modify the conventional 1T1R array by rotating the BLs by 90° , which is known as the pseudo-crossbar array architecture, as shown in Fig. 7(b). In weighted sum operation, all the transistors will be transparent when all WLs are turned on. Thus, the input vector voltages are provided to the BLs, and the weighted sum currents are read out through SLs in parallel. The weight update operation in pseudo-crossbar array is similar to that in crossbar array, as shown in Fig. 8. As the unselected WLs can turn off the transistors on unselected rows, no voltage bias is required for these unselected BLs thus pseudo-crossbar array can have save a lot of weight update energy compared to the crossbar array. In the simulator, the pseudo-crossbar array architecture can be designated by setting `cmosAccess=true` in **RealDevice** class of **Cell.cpp**.

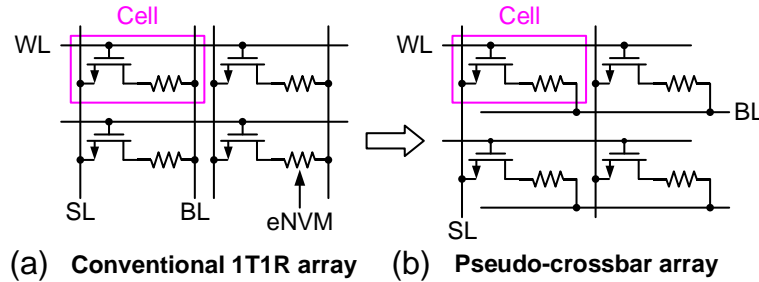


Fig. 7 Transformation from (a) conventional 1T1R array to (b) pseudo-crossbar array by 90° rotation of BL to enable weighted sum operation.




Fig. 8 Voltage bias scheme in the write operation of pseudo-crossbar array. Two separate phases for weight increase and decrease are required. In this example, the left cell of the selected cells will be updated in phase 1, while the right one will be updated in phase 2.

5.3 Array Peripheral Circuits

The periphery circuit modules used in both eNVM crossbar and pseudo-crossbar synaptic cores in Fig. 5 are described below:

1) Switch matrix

Switch matrixes are used for fully parallel voltage input to the array rows or columns. Fig. 9(a) shows the BL switch matrix for example. It consists of transmission gates that are connected to all the BLs, with control signals (B_1 to B_n) of the transmission gates stored in the registers (not shown here). In the weighted sum operation, the input vector signal is loaded to B_1 to B_n , which decide the BLs to be connected to either the read voltage or ground. In this way, the read voltage that is applied at the input of transmission gates can pass to the BLs and the weighted sums are read out through SLs in parallel. If the input vector is not 1 bit, it should be encoded using multiple clock cycles, as shown in Fig. 9(b). The reason why we do not use analog voltage to represent the input vector precision is the I-V nonlinearity of eNVM cell, which will cause the weighted sum distortion or inaccuracy. In the simulator, all the switch matrixes (`slSwitchMatrix`, `blSwitchMatrix` and `wlSwitchMatrix`) are instantiated from `SwitchMatrix` class in `SwitchMatrix.cpp`.

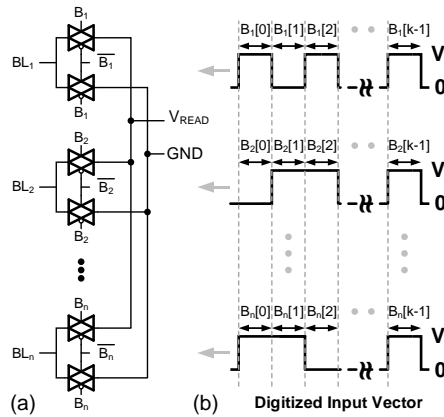


Fig. 9 (a) Transmission gates of the BL switch matrix in the weighted sum operation. A vector of control signals (B_1 to B_n) from the registers (not shown here) decide the BLs to be connected to either a voltage source or ground. (b) Control signals in a bit stream to represent the precision of the input vector.

2) Crossbar WL decoder

The crossbar WL decoder is modified from the traditional WL decoder. It has an additional feature to activate all the WLs for making all the transistors transparent for weighted sum. The crossbar WL decoder is constructed by attaching the follower circuits to every output row of the traditional decoder, as shown in Fig. 10. If $ALLOPEN=1$, the crossbar WL decoder will activate all the WLs no matter what input address is given, otherwise it will function as a traditional WL decoder. In the simulator, the crossbar WL decoder contains a traditional decoder (**wlDecoder**) instantiated from **RowDecoder** class in **RowDecoder.cpp** and a collection of follower circuits (**wlDecoderOutput**) instantiated from **WLDecoderOutput** class in **WLDecoderOutput.cpp**.

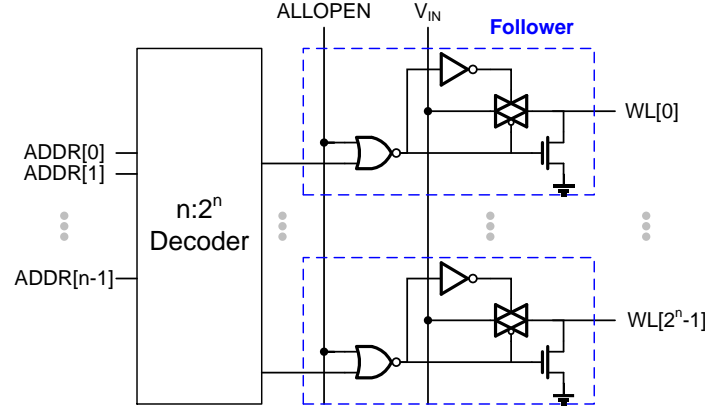


Fig. 10 Circuit diagram of the crossbar WL decoder. Follower circuit is attached to every row of the decoder to enable activation of all WLs when $ALLOPEN=1$.

3) Multiplexer (Mux) and Mux decoder

The Multiplexer (Mux) is used for sharing the read periphery circuits among synaptic array columns, because the array cell size is much smaller than the size of read periphery circuits and it will not be area-efficient to put all the read periphery circuits underneath the array. However, sharing the read periphery circuits among synaptic array columns inevitably increases the latency of weighted sum as time multiplexing is needed, which is controlled by the Mux decoder. In the simulator, the Mux (**mux**) is instantiated from **Mux** class in **Mux.cpp** and the Mux decoder (**muxDecoder**) is instantiated from **RowDecoder** class in **RowDecoder.cpp**.

4) Analog-to-digital read circuit

To convert these analog weighted sum currents to digital outputs, we use the read circuit in the following reference to employ the principle of the integrate-and-fire neuron model, as shown in Fig. 11(a). The read circuit integrates the weighted sum current on the finite capacitance of the array column. Once the voltage charges up above a certain threshold, the read circuit fires an output pulse and the capacitance is discharged back. The simulated waveform of integrated input voltage and the digital output spikes of the read circuit is shown in Fig. 11(b). The number of output spikes is proportional to the weight sum current. The precision required for this analog-to-digital conversion (ADC) determines the pulse width in each bit of the input vector. In the simulator, a collection of read circuits (**readCircuit**) is instantiated from **ReadCircuit** class in **ReadCircuit.cpp**.

D. Kadetotad, Z. Xu, A. Mohanty, P.-Y. Chen, B. Lin, J. Ye, S. Vruthula, S. Yu, Y. Cao, J.-S. Seo, "Parallel architecture with resistive crosspoint array for dictionary learning acceleration," *IEEE J. Emerg. Sel. Topics Circuits Syst. (JETCAS)*, vol. 5, no. 2, pp. 194-204, 2015.

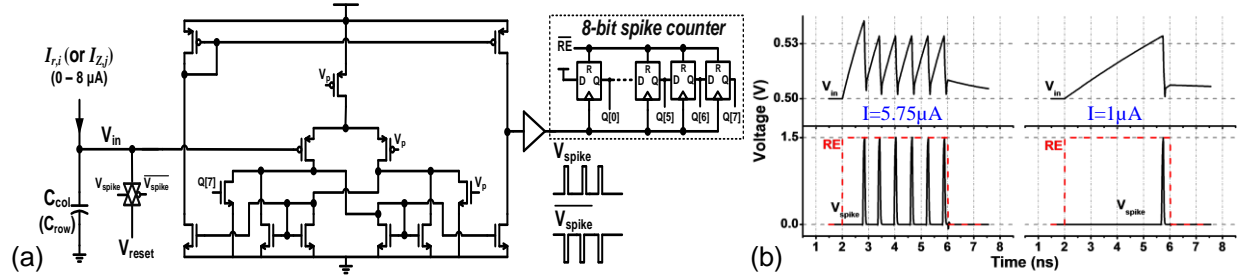


Fig. 11 (a) Design of a read circuit that employs the principle of the integrate-and-fire neuron model. (b) Simulated waveform of integrated input voltage and the digital output spikes of the read circuit.

5) Adder and shift register

The adder and shift register pair at the bottom of synaptic core performs shift and add of the weighted sum result at each input vector bit cycle (B_1 to B_n in Fig. 9(b)) to get the final weighted sum. The bit-width of the adder and shift register needs to be further extended depending on the precision of input vector. If the values in the input vector are only 1 bit, then the adder and shift register pair is not required. In the simulator, a collection of the adder and shift register pairs (**ShiftAdd**) is instantiated from **ShiftAdd** class in **ShiftAdd.cpp**, where **ShiftAdd** further contains a collection of adders (**adder**) instantiated from **Adder** class in **Adder.cpp** and a collection of registers (**dff**) instantiated from **DFF** class in **DFF.cpp**.

6. Algorithm Level: Multilayer Perceptron (MLP) Neural Network Architecture

At the algorithm level, we provide a simple 2-layer multilayer perceptron (MLP) neural network for performance benchmark. As shown in Fig. 12(a), the network consists of an input layer, hidden layer and output layer (the input layer is not included when counting the number of layers). MLP is a fully connected neural network, where each neuron node in one layer connects to every neuron node in the following layer. The connection between two neuron nodes is through a synapse with its strength representing the weight, where W_{IH} and W_{HO} are the weight matrix between input and hidden layer and between hidden and output layer, respectively. For the input image data, we post-processed the MNIST handwritten digits by cropping the edges of each image (making it 20×20 pixels). For the favor of hardware implementation, we also convert the images to black and white data in the simulator to reduce the design complexity of input encoding. By default, the network topology is 400(input layer)-100(hidden layer)-10(output layer). 400 neurons of input layer correspond to the 20×20 MNIST image, and 10 neurons of output layer correspond to 10 classes of digits. The users could change the network topology as needed. However, a new learning rate may be required to optimize the learning accuracy.

The learning applications that the network can implement in this simulator include the online learning and offline training with classification only. In online learning, the simulator emulates hardware to train the network with images randomly picked from the training dataset (60k images) and classify the testing dataset (10k images). In offline training with classification only, the network is pre-trained by software, and the MLP simulator only emulates hardware to classify the testing dataset. The training data file is **patch60000_train.txt** and its label file (the correct answers of the training data) is **label60000_train.txt**. The testing data file is **patch10000_test.txt** and its label file (the correct answers of the testing data) is **label60000_train.txt**.

The training process consists of two key operations, the feed forward (FF) and back propagation (BP). In feed forward, the input data are fed from the input layer and they will travel in a forward direction to the output layer via a series of weighted sum operation and neuron activation function along the way. The feed forward result at the output layer will then be compared with its correct answer (the label) to calculate its prediction error (the deviation). In back propagation, this error is propagated backward from the output layer to adjust the weights of each layer in a way that the prediction error is minimized. In this simulator, we use stochastic gradient decent method to update the weights in the back propagation. Different than the traditional gradient decent, the back propagation is performed after the feed forward of every image rather than that of the entire image dataset. On the other hand, the testing (classification) process only has the feed forward operation to make predictions. The weights in this process will not be changed.

Fig. 12(b) shows a schematic of a neuron node, which encapsulates the principles discussed above. The neuron takes the weighted sum result from its inward synapses and pass it through a 1-bit low-precision activation function. In this way, the offline classification, which is purely feed forward, can be realized in 1 bit. However, the computation on the back propagation of weight update generally needs higher precision to update the small errors, thus a high-precision activation function for the back propagation is still necessary.

Fig. 12(c) shows the circuit block diagram for hardware implementation of this 2-layer MLP neural network. The weighted sum operation is performed using the synaptic cores. However, the weights used in a regular synaptic array can only represent positive values ($W_H=0 \sim 1$), while the weights in algorithm can be either positive or negative values ($W_A=-1 \sim 1$). The algorithm's weighted sum is then expressed as

$$W_A V = (2W_H - J)V = 2W_H V - JV \quad (4)$$

where V is the input vector and J is the matrix of all ones that has the same dimension as W_A and W_H . In Eq. (4), $W_H V$ is the weighted sum output from the synaptic core. Therefore, we squeeze W_A from $(-1 \sim 1)$ to the range of W_H ($0 \sim 1$): i.e. -1 is mapped to 0 , 0 is mapped to 0.5 , and 1 is mapped to 1 . To reconstruct $W_A V$, we have to perform a two-step read from the array: first, we read out $W_H V$, and then multiply $W_H V$ by 2 using a 1 -bit left-shift, and then subtract JV (basically the sum of vector) from $W_H V$ through the adder at the periphery. The MSB (sign bit in 2 's complement notation) of the adder output will be the 1 -bit output of the low-precision activation function. It should be noted that we only consider the main sub-circuit modules for the neuron periphery at current stage of this simulator, and the hardware for BP error calculation as well as the detailed control logics will be included in the future release.

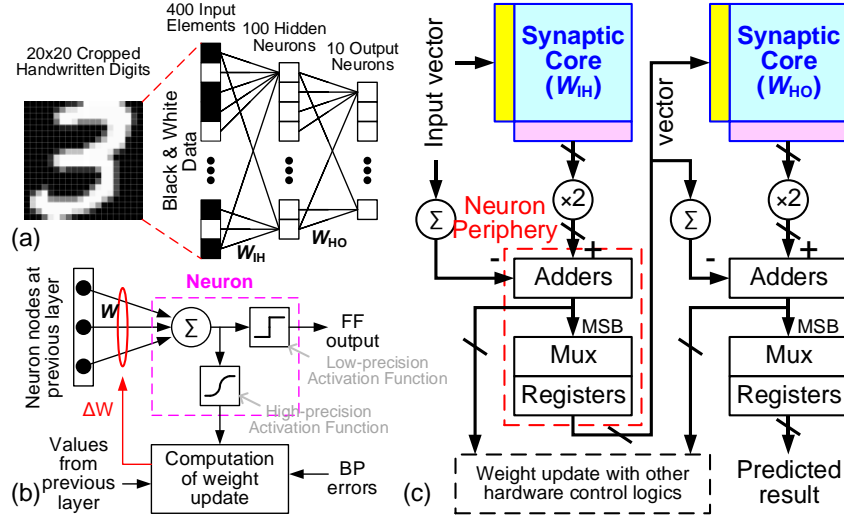


Fig. 12 (a) The 2-layer MLP neural network. (b) Schematic of a neuron node. (c) Circuit block diagram for hardware implementation of the 2-layer MLP network.

7. How to run MLP simulator (+NeuroSim)

Before running the simulator, the device-level and array-level parameters that the user may wish to modify are in **RealDevice** class of **Cell.cpp**, as shown in Fig. 13. The default values are for the Ag:Si example.

```

161  /* Real Device */
162  RealDevice::RealDevice(int x, int y) {
163      this->x = x; this->y = y; // Cell location: x (column) and y (row) start from index 0
164      maxConductance = 3.8462e-8; // Maximum cell conductance (S)
165      minConductance = 3.0769e-9; // Minimum cell conductance (S)
166      conductance = minConductance; // Current conductance (S) (dynamic variable)
167      readVoltage = 0.5; // On-chip read voltage (Vr) (V)
168      readPulseWidth = 5e-9; // Read pulse width (s) (will be determined by ADC)
169      writeVoltageLTP = 3.2; // Write voltage (V) for LTP or weight increase
170      writeVoltageLTD = 2.8; // Write voltage (V) for LTD or weight decrease
171      writePulseWidthLTP = 300e-6; // Write pulse width (s) for LTP or weight increase
172      writePulseWidthLTD = 300e-6; // Write pulse width (s) for LTD or weight decrease
173      writeEnergy = 0; // Dynamic variable for calculation of write energy (J)
174      maxNumLevelLTP = 97; // Maximum number of conductance states during LTP or weight increase
175      maxNumLevelLTD = 100; // Maximum number of conductance states during LTD or weight decrease
176      numPulse = 0; // Number of write pulses used in the most recent write operation (dynamic variable)
177      cmosAccess = true; // True: Pseudo-crossbar (1T1R), false: cross-point
178      resistanceAccess = 15e3; // The resistance of transistor (Ohm) in Pseudo-crossbar array when turned ON
179      nonlinearIV = false; // Consider I-V nonlinearity or not (Currently for cross-point array only)
180      NL = 10; // I-V nonlinearity in write scheme (the current ratio between Vw and Vw/2), assuming for the LTP side
181      if (nonlinearIV) { // Currently for cross-point array only
182          double Vr_exp = readVoltage; // XXX: Modify this value to Vr in the reported measurement data (can be different than readVoltage)
183          // Calculation of conductance at on-chip Vr
184          maxConductance = NonlinearConductance(maxConductance, NL, writeVoltageLTP, Vr_exp, readVoltage);
185          minConductance = NonlinearConductance(minConductance, NL, writeVoltageLTP, Vr_exp, readVoltage);
186      }
187      nonlinearWrite = true; // Consider weight update nonlinearity or not
188      readNoise = false; // Consider read noise or not
189      sigmaReadNoise = 0; // Sigma of read noise in gaussian distribution
190      gaussian_dist = new std::normal_distribution<double>(0, sigmaReadNoise); // Set up mean and stddev for read noise
191
192      std::mt19937 localGen; // It's OK not to use the external gen, since here the device-to-device variation is a one-time deal
193      localGen.seed(std::time(0));
194
195      /* Device-to-device weight update variation */
196      NL_LTP = 2.4; // LTP nonlinearity
197      NL_LTD = -4.88; // LTD nonlinearity
198      sigmaDtoD = 0; // Sigma of device-to-device weight update variation in gaussian distribution
199      gaussian_dist2 = new std::normal_distribution<double>(0, sigmaDtoD); // Set up mean and stddev for device-to-device weight update variation
200      paramALTP = getParamA(NL_LTP + (*gaussian_dist2)(localGen)) * maxNumLevelLTP; // Parameter A for LTP nonlinearity
201      paramALTLD = getParamA(NL_LTD + (*gaussian_dist2)(localGen)) * maxNumLevelLTD; // Parameter A for LTD nonlinearity
202
203      /* Cycle-to-cycle weight update variation */
204      sigmaCtoC = 0.035 * (maxConductance - minConductance); // Sigma of cycle-to-cycle weight update variation: defined as the percentage of conductance range
205      gaussian_dist3 = new std::normal_distribution<double>(0, sigmaCtoC); // Set up mean and stddev for cycle-to-cycle weight update variation
206
207      heightInFeatureSize = cmosAccess? 4 : 2; // Cell height = 4F (Pseudo-crossbar) or 2F (cross-point)
208      widthInFeatureSize = cmosAccess? 4 : 2; // Cell width = 4F (Pseudo-crossbar) or 2F (cross-point)
209  }

```

Fig. 13 Code snippet of **RealDevice** class in **Cell.cpp**

where **maxConductance** and **minConductance** are defined as $1/R_{on}$ and $1/R_{off}$ respectively, **readVoltage** and **readPulseWidth** are on-chip read voltage (V) and read pulse width (s). The specified value of **readPulseWidth** does not matter because it will be modified later by the read circuit module when it calculates the required pulse width for each integration cycle based on the ADC precision. **writeVoltageLTP** and **writePulseWidthLTP** are the write voltage (V) and the write pulse width (s) during LTP or weight increase. **writeVoltageLTD** and **writePulseWidthLTD** are also defined in the same way. **maxNumLevelLTP** and **maxNumLevelLTD** mean the maximum number of write pulse for LTP and LTD of the real device, which represents the precision. Besides, **NL_LTP** and **NL_LTD** represent the nonlinearity values for LTP and LTD, and **sigmaDtoD** and **sigmaCtoC** represent the device-to-device and cycle-to-cycle weight update variation, respectively. Currently the simulator only takes one value of the cycle-to-cycle weight update variation for **sigmaCtoC**. It is encouraged that the user selects the larger one in LTP and LTD for conservative estimation. In the example of Ag:a-Si, **NL_LTP** and **NL_LTD** are set to 2.40 and -4.88, which are obtained from the MATLAB fitting results. The cycle-to-cycle weight update variation is set to be 0.035 as the maximum variation of LTP and LTD is 3.5%.

Also, **cmosAccess** is used to choose the cell structure, or synaptic core type in other words. **cmosAccess=true** means the pseudo-crossbar (1T1R) array, while **cmosAccess=false** means the true crossbar array. If the cell is 1T1R, we need to define **resistanceAccess**, which is the turn-on resistance value of the transistor in 1T1R array. If the cell is crossbar, I-V nonlinearity **NL** can be specified as the current ratio between write voltage and half write voltage considering if a selector is added.

In the file **Param.cpp**, there are also network-level parameters that the user may wish to modify, as shown in Fig. 14.

```

42 Param::Param() {
43     /* MNIST dataset */
44     numMnistTrainImages = 60000; // # of training images in MNIST
45     numMnistTestImages = 10000; // # of testing images in MNIST
46
47     /* Algorithm parameters */
48     numTrainImagesPerEpoch = 8000; // # of training images per epoch
49     totalNumEpochs = 125; // Total number of epochs
50     interNumEpochs = 1; // Internal number of epochs (print out the results every interNumEpochs)
51     nInput = 400; // # of neurons in input layer
52     nHide = 100; // # of neurons in hidden layer
53     nOutput = 10; // # of neurons in output layer
54     alpha1 = 0.2; // Learning rate for the weights from input to hidden layer
55     alpha2 = 0.1; // Learning rate for the weights from hidden to output layer
56     maxWeight = 1; // Upper bound of weight value
57     minWeight = 0; // Lower bound of weight value
58
59     /* Hardware parameters */
60     useHardwareInTrainingFF = true; // Use hardware in the feed forward part of training or not (true: realistic hardware, false: ideal software)
61     useHardwareInTrainingWU = true; // Use hardware in the weight update part of training or not (true: realistic hardware, false: ideal software)
62     useHardwareInTraining = useHardwareInTrainingFF || useHardwareInTrainingWU; // Use hardware in the training or not
63     useHardwareInTestingFF = true; // Use hardware in the feed forward part of testing or not (true: realistic hardware, false: ideal software)
64     numBitInput = 1; // # of bits of the input data (1 for black and white data)
65     numBitPartialSum = 8; // # of bits of the digital output (partial weighted sum output)
66     pSumMaxHardware = pow(2, numBitPartialSum) - 1; // Max digital output value of partial weighted sum
67     numInputLevel = pow(2, numBitInput); // # of levels of the input data
68     numWeightBit = 6; // # of weight bits for pure algorithm (the # of weight bits for the device should be specified in maxNumLevelLTP and maxNumLevelLTD in Cell.cpp)
69     BWthreshold = 0.5; // The black and white threshold for numBitInput=1
70     Hthreshold = 0.5; // The spiking threshold for the hidden layer (dal in Train.cpp and Test.cpp)
71     numColMuxed = 16; // How many columns share 1 read circuit (for analog RRAM) or 1 S/A (for digital RRAM)
72     writeEnergyReport = true; // Report write energy calculation or not
73     NeuroSimDynamicPerformance = true; // Report the dynamic performance (latency and energy) in NeuroSim or not
74     relaxArrayCellHeight = 0; // True: relax the array cell height to standard logic cell height in the synaptic array
75     relaxArrayCellWidth = 0; // True: relax the array cell width to standard logic cell width in the synaptic array
76     arrayWireWidth = 40; // Array wire width (nm)
77     processNode = 14; // Technology node (nm)
78     clkFreq = 2e9; // Clock frequency (Hz)
79 }

```

Fig. 14 Code snippet of **Param.cpp**

where **numMnistTrainImages** and **numMnistTestImages** are the number of images in MNIST during training and testing respectively, **numTrainImagesPerEpoch** means the number of training images per epoch, while **interNumEpochs** represents the internal number of epochs within each printed epoch shown on the screen. In addition, **nInput**, **nHide** and **nOutput** are the number of neurons in input, hidden and output layers in the 2-layer MLP neural network, respectively.

The first four hardware parameters determine the learning configuration, which can be the following cases:

1. Online learning in hardware: **useHardwareInTrainingFF**, **useHardwareInTrainingWU** and **useHardwareInTestingFF** are all **true**
2. Offline learning in software and then classification only in hardware: **useHardwareInTrainingFF** and **useHardwareInTrainingWU** are **false**, while **useHardwareInTestingFF** is **true**
3. Pure learning in software: **useHardwareInTrainingFF**, **useHardwareInTrainingWU** and **useHardwareInTestingFF** are all **false**

For other hardware parameters, **numBitInput** means the number of bits of the input data. The hardware architecture design in this released version only allows **numBitInput=1** (black and white data), which should not be changed. **numBitPartialSum** represents the number of bits in the digital output (partial weighted sum output) of read circuit (ADC) at the periphery of the synaptic array. Also, **numWeightBit** means the number of weight bits for pure algorithm without consideration of hardware, and **numColMuxed** means the number of columns of the synaptic array sharing one read circuit in the array. Time-multiplexing is required if **numColMuxed** is greater than 1. For example, the total weighted sum latency will be increased by 16 times if **numColMuxed=16**.

Basically, the parameters for running the simulator are all included in the file **Cell.cpp** and **Param.cpp**. Whenever any change is made in the files, the codes has to be recompiled by using **make** command as stated in **Usage (Linux)** section. If the compilation is successful, the following screenshot of Fig. 15 can be expected:


```
[xpengl5@syu04 ~/MLP_simulator_with_NeuroSim]$ make
makefile:61: .depend: No such file or directory
g++ -c -fopenmp -O3 -std=c++0x -w Array.cpp -o Array.o
g++ -c -fopenmp -O3 -std=c++0x -w Cell.cpp -o Cell.o
g++ -c -fopenmp -O3 -std=c++0x -w IO.cpp -o IO.o
g++ -c -fopenmp -O3 -std=c++0x -w Mapping.cpp -o Mapping.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim.cpp -o NeuroSim.o
g++ -c -fopenmp -O3 -std=c++0x -w Param.cpp -o Param.o
g++ -c -fopenmp -O3 -std=c++0x -w Test.cpp -o Test.o
g++ -c -fopenmp -O3 -std=c++0x -w Train.cpp -o Train.o
g++ -c -fopenmp -O3 -std=c++0x -w formula.cpp -o formula.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/Adder.cpp -o NeuroSim/Adder.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/DFF.cpp -o NeuroSim/DFF.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/DecoderDriver.cpp -o NeuroSim/DecoderDriver.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/FunctionUnit.cpp -o NeuroSim/FunctionUnit.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/LUT.cpp -o NeuroSim/LUT.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/Mux.cpp -o NeuroSim/Mux.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/ReadCircuit.cpp -o NeuroSim/ReadCircuit.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/RowDecoder.cpp -o NeuroSim/RowDecoder.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/ShiftAdd.cpp -o NeuroSim/ShiftAdd.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/SubArray.cpp -o NeuroSim/SubArray.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/SwitchMatrix.cpp -o NeuroSim/SwitchMatrix.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/Technology.cpp -o NeuroSim/Technology.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/VoltageSenseAmp.cpp -o NeuroSim/VoltageSenseAmp.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/WLDecoderOutput.cpp -o NeuroSim/WLDecoderOutput.o
g++ -c -fopenmp -O3 -std=c++0x -w NeuroSim/formula.cpp -o NeuroSim/formula.o
g++ -c -fopenmp -O3 -std=c++0x -w main.cpp -o main.o
g++ -fopenmp -O3 -std=c++0x -w Array.o Cell.o IO.o Mapping.o NeuroSim.o Param.o Test.o Train.o formula.o NeuroSim/Adder.o NeuroSim/DFF.o NeuroSim/DecoderDriver.o NeuroSim/FunctionUnit.o NeuroSim/LUT.o NeuroSim/Mux.o NeuroSim/ReadCircuit.o NeuroSim/RowDecoder.o NeuroSim/ShiftAdd.o NeuroSim/SubArray.o NeuroSim/SwitchMatrix.o NeuroSim/Technology.o NeuroSim/VoltageSenseAmp.o NeuroSim/WLDecoderOutput.o NeuroSim/formula.o main.o -o main
```

Fig. 15 Output of make

Then, use `./main` or `make run` to run the program. For both commands, the program will print out the results at every epoch during the simulation. Compared to `./main`, `make run` can save the results to a log file with filename appended with the current time info, as shown in Fig. 16. With the default value of **totalNumEpochs=125** and **numTrainImagesPerEpoch=8000** for a total 1 million MNIST images, the simulation will approximately take about 40 mins with an Intel i7 CPU and 32 GB RAM.

```
[xpengl5@syu06 ~/MLP_simulator_with_NeuroSim]$ make run
stdbuf -o 0 ./main | tee log_20170619_155943.txt
Total SubArray (synaptic core) area=9.1425e-10 m^2
Total Neuron (neuron peripheries) area=1.5707e-10 m^2
Total area=1.0713e-09 m^2
Leakage power of subArrayIH is : 2.5285e-05 W
Leakage power of subArrayH0 is : 5.8906e-06 W
Leakage power of NeuronIH is : 3.5120e-06 W
Leakage power of NeuronH0 is : 6.0166e-07 W
Total leakage power of subArray is : 3.1176e-05 W
Total leakage power of Neuron is : 4.1137e-06 W

Accuracy at 125 epochs is : 71.22%
Read latency=2.9897e-02 s
Write latency=3.0000e+07 s
Read energy=1.4714e-04 J
Write energy=1.0314e+02 J
```

Fig. 16 Output of the simulation

At the end of simulation, it is expected to have similar results for the Ag:a-Si example in Table 2:

Accuracy	Area (m ²)	Read Latency (s)	Write Latency (s)	Read Energy (J)	Write Energy (J)
71.22%	1.0713e ⁻⁹	2.9897e ⁻²	3.0000e ⁷	1.4714e ⁻⁴	1.0314e ²

Table 2 Final results of learning accuracy and circuit-level performance

For the accuracy of pure software learning as the reference, the users could change the learning modes in **Param.cpp** as discussed above, and it is 96~97% for a network topology 400-100-10.